

11: Machine Learning System Design

[Previous](#) [Next](#) [Index](#)

Machine learning systems design

- In this section we'll touch on how to put together a system
- Previous sections have looked at a wide range of different issues in significant focus
- This section is less mathematical, but material will be very useful non-the-less
 - Consider the system approach
 - You can understand all the algorithms, but if you don't understand how to make them work in a complete system that's no good!

Prioritizing what to work on - spam classification example

- The idea of prioritizing what to work on is perhaps the most important skill programmers typically need to develop
 - It's so easy to have many ideas you want to work on, and as a result do none of them well, because doing one well is harder than doing six superficially
 - So you need to make sure you complete projects
 - Get something "shipped" - even if it doesn't have all the bells and whistles, that final 20% getting it ready is often the toughest
 - If you only release when you're totally happy you rarely get practice doing that final 20%
 - So, back to machine learning...
- Building a spam classifier
- Spam is email advertising

From: cheapsales@buystufffromme.com
To: ang@cs.stanford.edu
Subject: Buy now!

Deal of the week! Buy now!
Rolex w4tchs - \$100
Medicine (any kind) - \$50
Also low cost M0rgages available.

Spam

From: Alfred Ng
To: ang@cs.stanford.edu
Subject: Christmas dates?

Hey Andrew,
Was talking to Mom about plans for Xmas. When do you get off work. Meet Dec 22?
Alf

Non-spam

- What kind of features might we define
 - Spam (1)
 - Misspelled word
 - Not spam (0)
 - Real content
- How do we build a classifier to distinguish between the two
 - Feature representation
 - How do represent x (features of the email)?
 - $y = \text{spam (1) or not spam (0)}$

One approach - choosing your own features

- Chose 100 words which are indicative of an email being spam or not spam
 - Spam --> e.g. buy, discount, deal
 - Non spam --> Andrew, now
 - All these words go into one long vector
- Encode this into a **reference vector**
 - See which words appear in a message
- Define a feature vector x
 - Which is 0 or 1 if a word corresponding word in the reference vector is present or not
 - This is a bitmap of the word content of your email
 - i.e. don't recount if a word appears more than once

$$x = \begin{bmatrix} 0 \\ 1 \\ 1 \\ 0 \\ \vdots \\ 1 \\ \vdots \end{bmatrix} \begin{matrix} \text{andrew} \\ \text{buy} \\ \text{deal} \\ \text{discount} \\ \vdots \\ \text{now} \\ \vdots \end{matrix}$$

- In practice its more common to have a training set and pick the most frequently n words, where n is 10 000 to 50 000
 - So here you're not specifically choosing your own features, but you are choosing *how* you select them from the training set data

What's the best use of your time to improve system accuracy?

- Natural inclination is to collect lots of data
 - Honey pot anti-spam projects try and get fake email addresses into spammers' hands, collect loads of spam
 - This doesn't always help though
- Develop sophisticated features based on email routing information (contained in email header)
 - Spammers often try and obscure origins of email
 - Send through unusual routes
- Develop sophisticated features for message body analysis
 - Discount == discounts?
 - DEAL == deal?
- Develop sophisticated algorithm to detect misspelling
 - Spammers use misspelled word to get around detection systems
- Often a research group **randomly focus on one option**
 - May not be the most fruitful way to spend your time
 - If you brainstorm a set of options this is **really good**
 - Very tempting to just try something

Error analysis

- When faced with a ML problem lots of ideas of how to improve a problem
 - Talk about error analysis - how to better make decisions
- If you're building a machine learning system often good to start by building a simple algorithm which you can implement quickly
 - Spend at most 24 hours developing an initially bootstrapped algorithm
 - Implement and test on cross validation data

- Plot learning curves to decide if more data, features etc will help algorithmic optimization
 - Hard to tell in advance what is important
 - Learning curves really help with this
 - Way of avoiding **premature optimization**
 - We should let evidence guide decision making regarding development trajectory
- **Error analysis**
 - Manually examine the samples (in cross validation set) that your algorithm made errors on
 - See if you can work out why
 - Systematic patterns - help design new features to avoid these shortcomings
 - e.g.
 - Built a spam classifier with 500 examples in CV set
 - Here, error rate is high - gets 100 wrong
 - Manually look at 100 and categorize them depending on features
 - e.g. type of email
 - Looking at those email
 - May find **most common type** of spam emails are pharmacy emails, phishing emails
 - See which type is most common - focus your work on those ones
 - What **features would have helped** classify them correctly
 - e.g. deliberate misspelling
 - Unusual email routing
 - Unusual punctuation
 - May find some "spammer technique" is causing a lot of your misses
 - Guide a way around it
- Importance of **numerical evaluation**
 - Have a way of numerically evaluated the algorithm
 - If you're developing an algorithm, it's really good to have some performance calculation which gives a single real number to tell you how well its doing
 - e.g.
 - Say were deciding if we should treat a set of similar words as the same word
 - This is done by stemming in NLP (e.g. "Porter stemmer" looks at the etymological stem of a word)
 - This may make your algorithm better or worse
 - Also worth consider weighting error (false positive vs. false negative)
 - e.g. is a false positive really bad, or is it worth have a few of one to improve performance a lot
 - Can use numerical evaluation to compare the changes
 - See if a change improves an algorithm or not
 - A single real number may be hard/complicated to compute
 - But makes it much easier to evaluate how changes impact your algorithm
- You should do error analysis on the cross validation set instead of the test set

Error metrics for skewed analysis

- Once case where it's hard to come up with good error metric - skewed classes
- Example
 - Cancer classification
 - Train logistic regression model $h_\theta(x)$ where
 - Cancer means $y = 1$
 - Otherwise $y = 0$
 - Test classifier on test set
 - Get 1% error
 - So this looks pretty good..
 - But only 0.5% have cancer

- Now, 1% error looks very bad!
- So when one number of examples is very small this is an example of skewed classes
 - LOTS more of one class than another
 - So standard error metrics aren't so good
- Another example
 - Algorithm has 99.2% accuracy
 - Make a change, now get 99.5% accuracy
 - Does this really represent an improvement to the algorithm?
 - Did we do something useful, or did we just create something which predicts $y = 0$ more often
 - Get very low error, but classifier is still not great


Precision and recall

- Two new metrics - **precision** and **recall**
 - Both give a value between 0 and 1
 - Evaluating classifier on a test set
 - For a test set, the actual class is 1 or 0
 - Algorithm predicts some value for class, predicting a value for each example in the test set
 - Considering this, classification can be
 - True positive (we guessed 1, it was 1)
 - False positive (we guessed 1, it was 0)
 - True negative (we guessed 0, it was 0)
 - False negative (we guessed 0, it was 1)
 - **Precision**
 - *How often does our algorithm cause a false alarm?*
 - Of all patients we predicted have cancer, what fraction of them *actually* have cancer
 - = true positives / # predicted positive
 - = true positives / (true positive + false positive)
 - High precision is good (i.e. closer to 1)
 - You want a big number, because you want false positive to be as close to 0 as possible
 - **Recall**
 - *How sensitive is our algorithm?*
 - Of all patients in set that actually have cancer, what fraction did we correctly detect
 - = true positives / # actual positives
 - = true positive / (true positive + false negative)
 - High recall is good (i.e. closer to 1)
 - You want a big number, because you want false negative to be as close to 0 as possible
 - By computing precision and recall get a better sense of how an algorithm is doing
 - This can't really be gamed
 - Means we're much more sure that an algorithm is good
 - Typically we say the presence of a rare class is what we're trying to determine (e.g. positive (1) is the existence of the rare thing)

Trading off precision and recall

- For many applications we want to control the trade-off between precision and recall
- Example
 - Trained a logistic regression classifier
 - Predict 1 if $h_{\theta}(x) \geq 0.5$
 - Predict 0 if $h_{\theta}(x) < 0.5$
 - This classifier may give some value for precision and some value for recall
 - Predict 1 only if very confident
 - One way to do this modify the algorithm we could modify the prediction threshold

- Predict 1 if $h_{\theta}(x) \geq 0.8$
 - Predict 0 if $h_{\theta}(x) < 0.2$
 - Now we can be more confident a 1 is a true positive
 - But classifier has lower recall - predict $y = 1$ for a smaller number of patients
 - Risk of false negatives
- Another example - avoid false negatives
 - This is probably worse for the cancer example
 - Now we may set to a lower threshold
 - Predict 1 if $h_{\theta}(x) \geq 0.3$
 - Predict 0 if $h_{\theta}(x) < 0.7$
 - i.e. 30% chance they have cancer
 - So now we have a higher recall, but lower precision
 - Risk of false positives, because we're less discriminating in deciding what means the person has cancer
- This threshold defines the trade-off
 - We can show this graphically by plotting precision vs. recall



- This curve can take many different shapes depending on classifier details
- Is there a way to automatically choose the threshold
 - Or, if we have a few algorithms, how do we compare different algorithms or parameter sets?

	Precision(P)	Recall (R)
Algorithm 1	0.5	0.4
Algorithm 2	0.7	0.1
Algorithm 3	0.02	1.0

- How do we decide which of these algorithms is best?
 - We spoke previously about using a single real number evaluation metric
 - By switching to precision/recall we have two numbers
 - Now comparison becomes harder
 - Better to have just one number
 - How can we convert P & R into one number?
 - One option is the average - $(P + R)/2$
 - This is not such a good solution
 - Means if we have a classifier which predicts $y = 1$ all the time you get a high recall and low precision
 - Similarly, if we predict Y rarely get high precision and low recall
 - So averages here would be 0.45, 0.4 and 0.51
 - 0.51 is best, despite having a recall of 1 - i.e. predict $y=1$ for everything
 - So average isn't great
 - **F₁Score (fscore)**
 - $= 2 * (PR / [P + R])$


- Fscore is like taking the average of precision and recall giving a higher weight to the lower value
- Many formulas for computing comparable precision/accuracy values
 - If $P = 0$ or $R = 0$ the Fscore = 0
 - If $P = 1$ and $R = 1$ then Fscore = 1
 - The remaining values lie between 0 and 1
- Threshold offers a way to control trade-off between precision and recall
- Fscore gives a single real number evaluation metric
 - If you're trying to automatically set the threshold, one way is to try a range of threshold values and evaluate them on your cross validation set
 - Then pick the threshold which gives the best fscore.

Data for machine learning

- Now switch tracks and look at how much data to train on
- On early videos caution on just blindly getting more data
 - Turns out under certain conditions getting more data is a very effective way to improve performance

Designing a high accuracy learning system

- There have been studies of using different algorithms on data
 - Data - confusing words (e.g. two, to or too)
 - Algorithms
 - Perceptron (logistic regression)
 - Winnow
 - Like logistic regression
 - Used less now
 - Memory based
 - Used less now
 - Talk about this later
 - Naive Bayes
 - Cover later
 - Varied training set size and tried algorithms on a range of sizes



- What can we conclude
 - Algorithms give remarkably similar performance
 - As training set sizes increases accuracy increases
 - Take an algorithm, give it more data, should beat a "better" one with less data
 - Shows that
 - Algorithm choice is pretty similar
 - More data helps

- When is this true and when is it not?
 - If we can correctly assume that features x have enough information to predict y accurately, then more data will probably help
 - A useful test to determine if this is true can be, "given x , can a human expert predict y ?"
 - So lets say we use a learning algorithm with many parameters such as logistic regression or linear regression with many features, or neural networks with many hidden features
 - These are powerful learning algorithms with many parameters which can fit complex functions
 - Such algorithms are low bias algorithms
 - Little systemic bias in their description - flexible
 - Use a small training set
 - Training error should be small
 - Use a very large training set
 - If the training set error is close to the test set error
 - Unlikely to over fit with our complex algorithms
 - So the test set error should also be small
 - Another way to think about this is we want our algorithm to have low bias and low variance
 - Low bias --> use complex algorithm
 - Low variance --> use large training set